

NITIN DOODHIYA

+91 7753938247 — nitindoodhiya@gmail.com — LinkedIn: [nitin-98](#) — GitHub: [nitindoodhiya](#) — Hyderabad, India

SUMMARY

Senior Software Engineer with 5+ years of experience designing and owning scalable distributed systems and GenAI-powered platforms. Proven track record architecting high-throughput pipelines (1M+ assets/day), building agentic AI frameworks (MCP Gateway, LLM-driven schema translation), saving \$90k/year in infrastructure costs, and delivering production services with sub-200ms latency. Experienced mentoring engineers, driving technical design, and partnering with product/business stakeholders to ship enterprise AI products.

TECHNICAL SKILLS

Languages: Java, Python, JavaScript, C++

Backend: Spring Boot, FastAPI, Node.js, REST/GraphQL APIs, Kafka, Docker, Microservices, Linux

Cloud & Data: AWS (Lambda, ECS, Fargate, S3), GCP, Azure, MySQL, MongoDB, Redis, Elasticsearch

Core: Distributed Systems, System Design, Event-Driven Architecture, API Design, Generative & Agentic AI, Vector Search, LLM Integration (Claude, OpenAI), MCP, Observability

EXPERIENCE

Software Engineer 2 — TypefaceAI, Hyderabad

Feb 2024 – Present

Enterprise Generative AI Content Platform

- Slashed \$90k/year in infrastructure and API costs by architecting a custom OpenAPI and MCP Gateway using LLM-driven schema translation for enterprise clients.
- Designed and built a configuration-driven Connector Service in Java publishing 500k assets/day across 40+ external platforms; reduced operational effort by 90% and user actions by 80%.
- Architected a Spring Batch orchestration engine managing complex workflows with high execution reliability and fault tolerance under production load.
- Delivered Typeface's official DAM integration plugin for Salesforce, driving adoption across Salesforce's enterprise ecosystem.
- Engineered CDN-backed public asset hosting layer (<200 ms cross-platform latency) powering production asset hosting for tier-1 enterprise clients, including one of the world's largest banks.
- Spearheaded technical direction for high-scale asset discovery and indexing on Elasticsearch, processing 1M+ daily assets for vector search while mentoring team members on production standards.

Backend Engineer — Rizzle, Hyderabad

Jun 2022 – Feb 2024

Video & Content Platform

- Designed and owned distributed video processing pipeline on AWS Lambda + Fargate; accelerated video generation speed by 30% while improving scalability and cost efficiency.
- Optimized event-driven streaming pipelines using Apache Kafka, truncating email delivery latency by 40% and processing time by 30%; defined topic partition design, consumer groups, and DLQ patterns.
- Engineered core backend infrastructure (rate limiting, exponential retries, distributed key generation) that improved system reliability under high concurrent load; led production incident response.
- Integrated Stripe payment system with <2 s latency and optimized Elasticsearch + S3 data retrieval paths for production query performance.

Software Engineer — Citi, Pune

Jul 2021 – Jun 2022

Global financial services company

- Spearheaded template management platform supporting 100+ templates across communication channels; implemented custom AES-based secure transfer protocols ensuring data integrity across teams.

GenAI / Agentic PROJECTS

Daichi (Python, Claude Hooks)

Built always-on Mac assistant managing multiple concurrent Claude sessions with real-time authorization, notifications, and background routing; demonstrated multi-agent orchestration patterns.

Cerebrum (Python, Claude)

Designed centralized orchestrator enabling multiple Claude sessions to share knowledge across repositories without re-reading codebases; reduced token consumption by ~90%.

Flatmate Agent (FastAPI, LLM)

Built conversational LLM agent with structured data extraction and memory for housing search workflows.

EDUCATION

B.Tech in Computer Science — Motilal Nehru National Institute of Technology (MNNIT), Allahabad — 2021